

For Credit card dataset perform the following

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm
```

```
In [2]: credit_df = pd.read_csv("credit_dataset.csv")
credit_df
```

```
Out[2]:
```


	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed	I
0	1	30.83	0.000	1	1	Industrials	White	1.25	
1	0	58.67	4.460	1	1	Materials	Black	3.04	
2	0	24.50	0.500	1	1	Materials	Black	1.50	
3	1	27.83	1.540	1	1	Industrials	White	3.75	
4	1	20.17	5.625	1	1	Industrials	White	1.71	
...
685	1	21.08	10.085	0	0	Education	Black	1.25	
686	0	22.67	0.750	1	1	Energy	White	2.00	
687	0	25.25	13.500	0	0	Healthcare	Latino	2.00	
688	1	17.92	0.205	1	1	ConsumerStaples	White	0.04	
689	1	35.00	3.375	1	1	Energy	Black	8.29	

690 rows × 16 columns

1.spot outliers in Income using bivariate plot

```
In [7]: credit_df.plot('Age', 'Income', kind='scatter', marker='o')
```


```
Out[7]: <AxesSubplot:xlabel='Age', ylabel='Income'>
```



2. Spot outliers in any one feature using box plot

```
In [8]: credit_df['Age'].plot(kind='box')
```


```
Out[8]: <AxesSubplot:>
```



3.spot outliers using histogram plot

In [9]: `credit_df['Debt'].plot(kind='hist')`


Out[9]: `<AxesSubplot:ylabel='Frequency'>`



4.Detect outliers in any one feature using IQR method

```
In [10]: import numpy as np
import matplotlib.pyplot as plt
np.random.seed(102)
grades = np.concatenate([[50,52,53,55,56,60,61,62,65,67]*20,
np.random.randint(0, 101, size=300)])
Q1 = np.percentile(grades , 25)
Q3 = np.percentile(grades , 75)
Q1,Q3 = np.percentile(grades , [25,75])
IQR = Q3 - Q1
ul = Q3+1.5*IQR
ll = Q1-1.5*IQR
outliers = grades[(grades > ul) | (grades < ll)]
print(outliers)
fig = plt.figure(figsize=(6,5))
hypo = np.random.randint(20, 81, size=500)
plt.boxplot([grades, hypo], widths=0.5)
plt.xticks([1,2],['Our data', 'Hypothetical data'])
plt.ylabel('Grade')
plt.title('Box plot of midterm grade')
plt.show()
```

```
[ 0  7  4  3  0  4  2  7  6 100  1  3  0  3 100 100 100 100
  4  0  3  6  6  6 100  7  6 100 100  6  3  6  1  6  0]
```



5. Detect outliers using z-score method

```
In [12]: import numpy as np
data = [1, 2, 2, 2, 3, 1, 1, 15, 2, 2, 2, 3, 1, 1, 2]
mean = np.mean(data)
std = np.std(data)
print('mean of the dataset is', mean)
print('std. deviation is', std)
threshold = 3
outlier = []
for i in data:
    z = (i-mean)/std
    if z > threshold:
        outlier.append(i)
print('outlier in dataset of Z score is', outlier)
```

```
mean of the dataset is 2.6666666666666665
std. deviation is 3.3598941782277745
outlier in dataset of Z score is [15]
```

6. Treat outliers by Deleting observations

```
In [20]: q1 = credit_df["Age"].quantile(0.25)
q3 = credit_df["Age"].quantile(0.75)
iqr = q3-q1
upper_bound = q3+(1.5*iqr)
lower_bound = q1-(1.5*iqr)
```

```
In [21]: upperIndex = credit_df[credit_df["Age"]>upper_bound].index
credit_df.drop(upperIndex,inplace=True)
lowerIndex = credit_df[credit_df["Age"]<lower_bound].index
```

```
credit_df.drop(lowerIndex,inplace=True)
credit_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 553 entries, 0 to 689
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                553 non-null   int64
1   Age                   553 non-null   float64
2   Debt                  553 non-null   float64
3   Married               553 non-null   int64
4   BankCustomer          553 non-null   int64
5   Industry              553 non-null   object
6   Ethnicity             553 non-null   object
7   YearsEmployed         553 non-null   float64
8   PriorDefault          553 non-null   int64
9   Employed              553 non-null   int64
10  CreditScore           553 non-null   int64
11  DriversLicense        553 non-null   int64
12  Citizen               553 non-null   object
13  ZipCode               553 non-null   int64
14  Income                553 non-null   int64
15  Approved              553 non-null   int64
dtypes: float64(3), int64(10), object(3)
memory usage: 73.4+ KB
```

7.Treat outliers using imputations

imputations using mean

```
In [22]: m = np.mean(credit_df['Age'])
print('mean:',m)
for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        titanic_df['Age'] = titanic_df['Age'].replace(i,m)
```

mean: 29.347486437613018

imputations using median

```
In [24]: m = credit_df['Age'].median()
print("median",m)
for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,m)
```

median 27.58

imputations using zero

```
In [25]: for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,0)
```

Univariate, Bivariate and Multivariate Analysis

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import math
```

```
In [3]: card_approval_df=pd.read_csv('clean_dataset.csv')
print(card_approval_df.head())
```

	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	\
0	1	30.83	0.000	1	1	Industrials	White	
1	0	58.67	4.460	1	1	Materials	Black	
2	0	24.50	0.500	1	1	Materials	Black	
3	1	27.83	1.540	1	1	Industrials	White	
4	1	20.17	5.625	1	1	Industrials	White	

	YearsEmployed	PriorDefault	Employed	CreditScore	DriversLicense	\
0	1.25	1	1	1	0	
1	3.04	1	1	6	0	
2	1.50	1	0	0	0	
3	3.75	1	1	5	1	
4	1.71	1	0	0	0	

	Citizen	ZipCode	Income	Approved
0	ByBirth	202	0	1
1	ByBirth	43	560	1
2	ByBirth	280	824	1
3	ByBirth	100	3	1
4	ByOtherMeans	120	0	1

```
In [4]: print(card_approval_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Gender                690 non-null    int64
1   Age                   690 non-null    float64
2   Debt                  690 non-null    float64
3   Married               690 non-null    int64
4   BankCustomer          690 non-null    int64
5   Industry              690 non-null    object
6   Ethnicity             690 non-null    object
7   YearsEmployed         690 non-null    float64
8   PriorDefault          690 non-null    int64
9   Employed              690 non-null    int64
10  CreditScore           690 non-null    int64
11  DriversLicense        690 non-null    int64
12  Citizen               690 non-null    object
13  ZipCode               690 non-null    int64
14  Income                690 non-null    int64
15  Approved              690 non-null    int64
dtypes: float64(3), int64(10), object(3)
memory usage: 86.4+ KB
None
```

```
In [5]: card_approval_df.duplicated().sum()

Out[5]: 0
```

Univariate Analysis of continuous Variables


```
In [6]: card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']].describe()
```

Out[6]:

	Age	Debt	YearsEmployed	CreditScore	Income
count	690.000000	690.000000	690.000000	690.000000	690.000000
mean	31.514116	4.758725	2.223406	2.400000	1017.385507
std	11.860245	4.978163	3.346513	4.86294	5210.102598
min	13.750000	0.000000	0.000000	0.000000	0.000000
25%	22.670000	1.000000	0.165000	0.000000	0.000000
50%	28.460000	2.750000	1.000000	0.000000	5.000000
75%	37.707500	7.207500	2.625000	3.000000	395.500000
max	80.250000	28.000000	28.500000	67.000000	100000.000000

```
In [7]: sns.histplot(card_approval_df.Age, kde=True)

Out[7]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```




Univariate Analysis of categorical Variables

```
In [8]: sns.countplot(card_approval_df.Gender)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[8]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```





```
In [9]: sns.countplot(card_approval_df.Ethnicity)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(  
<AxesSubplot:xlabel='Ethnicity', ylabel='count'>
```

Out[9]:



Bivariate analysis of continuous variable

```
In [10]: card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']].corr()
```

```
Out[10]:
```


	Age	Debt	YearsEmployed	CreditScore	Income
Age	1.000000	0.202177	0.391464	0.187327	0.018719
Debt	0.202177	1.000000	0.298902	0.271207	0.123121
YearsEmployed	0.391464	0.298902	1.000000	0.322330	0.051345
CreditScore	0.187327	0.271207	0.322330	1.000000	0.063692
Income	0.018719	0.123121	0.051345	0.063692	1.000000

```
In [12]: sns.scatterplot(card_approval_df.YearsEmployed, card_approval_df.Income)
plt.ylim(0, 20000)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
(0.0, 20000.0))
```

```
Out[12]:
```



Bivariate Analysis of Categorical Variables vs Continuous Variables


```
In [13]: card_approval_df.groupby(by='Approved').agg('mean')[['Age', 'Debt', 'YearsEmployed',
```

Out[13]:

	Age	Debt	YearsEmployed	CreditScore	Income
Approved					
0	29.773029	3.839948	1.257924	0.631854	198.605744
1	33.686221	5.904951	3.427899	4.605863	2038.859935

```
In [19]: plt.figure(figsize=(5,4))  
sns.kdeplot(data=card_approval_df,x='Age',hue='Approved',fill=True)
```


Out[19]: <AxesSubplot:xlabel='Age', ylabel='Density'>



Bivariate Analysis of Categorical Variables vs Categorical Variables

```
In [20]: sns.countplot(data=card_approval_df,x='Approved',hue='Gender')
```


Out[20]: <AxesSubplot:xlabel='Approved', ylabel='count'>



Multivariate Analysis

```
In [21]: sns.pairplot(data=card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']])
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x23b342f37c0>
```



In []:

For Credit card dataset perform the following

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm
```

```
In [2]: credit_df = pd.read_csv("credit_dataset.csv")
credit_df
```

```
Out[2]:
```


	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	YearsEmployed	I
0	1	30.83	0.000	1	1	Industrials	White	1.25	
1	0	58.67	4.460	1	1	Materials	Black	3.04	
2	0	24.50	0.500	1	1	Materials	Black	1.50	
3	1	27.83	1.540	1	1	Industrials	White	3.75	
4	1	20.17	5.625	1	1	Industrials	White	1.71	
...
685	1	21.08	10.085	0	0	Education	Black	1.25	
686	0	22.67	0.750	1	1	Energy	White	2.00	
687	0	25.25	13.500	0	0	Healthcare	Latino	2.00	
688	1	17.92	0.205	1	1	ConsumerStaples	White	0.04	
689	1	35.00	3.375	1	1	Energy	Black	8.29	

690 rows × 16 columns

1.spot outliers in Income using bivariate plot

```
In [7]: credit_df.plot('Age', 'Income', kind='scatter', marker='o')
```


```
Out[7]: <AxesSubplot:xlabel='Age', ylabel='Income'>
```



2. Spot outliers in any one feature using box plot

```
In [8]: credit_df['Age'].plot(kind='box')
```


```
Out[8]: <AxesSubplot:>
```



3.spot outliers using histogram plot

In [9]: `credit_df['Debt'].plot(kind='hist')`

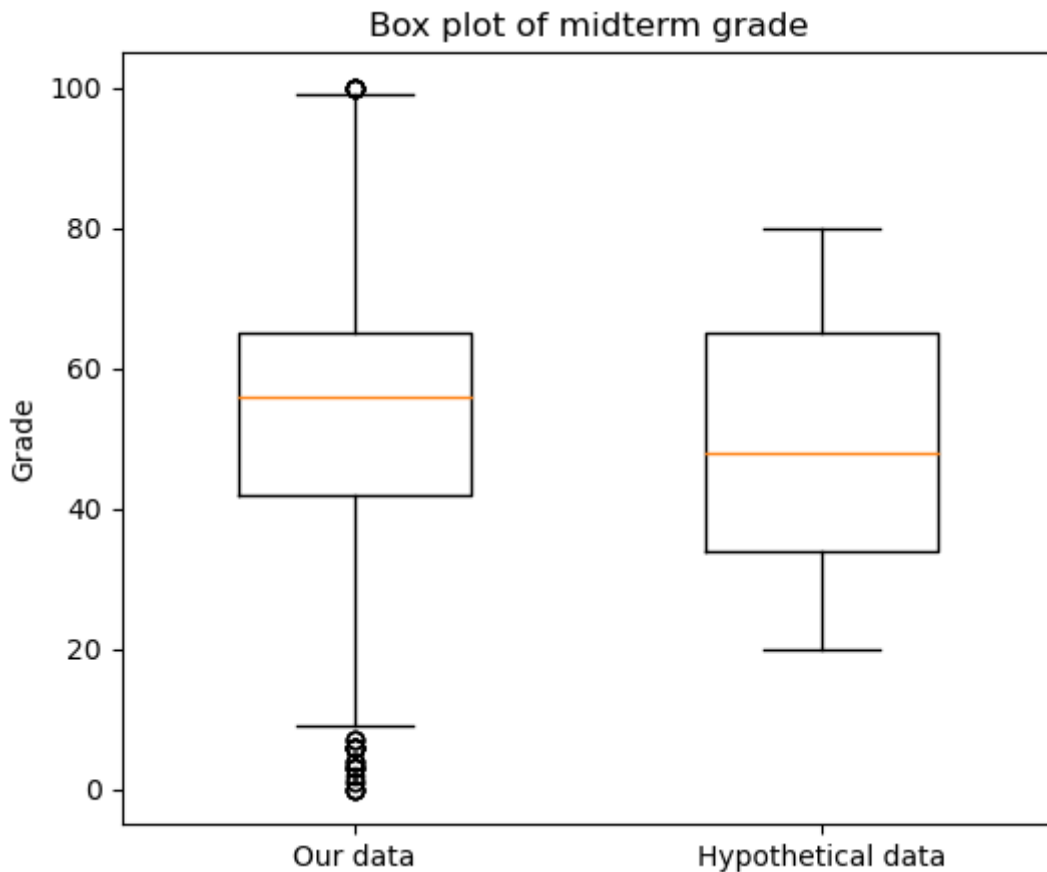
Out[9]: `<AxesSubplot:ylabel='Frequency'>`



4.Detect outliers in any one feature using IQR method

```
In [10]: import numpy as np
import matplotlib.pyplot as plt
np.random.seed(102)
grades = np.concatenate([[50,52,53,55,56,60,61,62,65,67]*20,
np.random.randint(0, 101, size=300)])
Q1 = np.percentile(grades , 25)
Q3 = np.percentile(grades , 75)
Q1,Q3 = np.percentile(grades , [25,75])
IQR = Q3 - Q1
ul = Q3+1.5*IQR
ll = Q1-1.5*IQR
outliers = grades[(grades > ul) | (grades < ll)]
print(outliers)
fig = plt.figure(figsize=(6,5))
hypo = np.random.randint(20, 81, size=500)
plt.boxplot([grades, hypo], widths=0.5)
plt.xticks([1,2],['Our data', 'Hypothetical data'])
plt.ylabel('Grade')
plt.title('Box plot of midterm grade')
plt.show()
```

```
[ 0  7  4  3  0  4  2  7  6 100  1  3  0  3 100 100 100 100
  4  0  3  6  6  6 100  7  6 100 100  6  3  6  1  6  0]
```

5. Detect outliers using z-score method

```
In [12]: import numpy as np
data = [1, 2, 2, 2, 3, 1, 1, 15, 2, 2, 2, 3, 1, 1, 2]
mean = np.mean(data)
std = np.std(data)
print('mean of the dataset is', mean)
print('std. deviation is', std)
threshold = 3
outlier = []
for i in data:
    z = (i-mean)/std
    if z > threshold:
        outlier.append(i)
print('outlier in dataset of Z score is', outlier)
```

```
mean of the dataset is 2.6666666666666665
std. deviation is 3.3598941782277745
outlier in dataset of Z score is [15]
```

6. Treat outliers by Deleting observations

```
In [20]: q1 = credit_df["Age"].quantile(0.25)
q3 = credit_df["Age"].quantile(0.75)
iqr = q3-q1
upper_bound = q3+(1.5*iqr)
lower_bound = q1-(1.5*iqr)
```

```
In [21]: upperIndex = credit_df[credit_df["Age"]>upper_bound].index
credit_df.drop(upperIndex,inplace=True)
lowerIndex = credit_df[credit_df["Age"]<lower_bound].index
```

```
credit_df.drop(lowerIndex,inplace=True)
credit_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 553 entries, 0 to 689
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                553 non-null   int64
1   Age                   553 non-null   float64
2   Debt                  553 non-null   float64
3   Married               553 non-null   int64
4   BankCustomer          553 non-null   int64
5   Industry              553 non-null   object
6   Ethnicity             553 non-null   object
7   YearsEmployed         553 non-null   float64
8   PriorDefault          553 non-null   int64
9   Employed              553 non-null   int64
10  CreditScore           553 non-null   int64
11  DriversLicense        553 non-null   int64
12  Citizen               553 non-null   object
13  ZipCode               553 non-null   int64
14  Income                553 non-null   int64
15  Approved              553 non-null   int64
dtypes: float64(3), int64(10), object(3)
memory usage: 73.4+ KB
```

7.Treat outliers using imputations

imputations using mean

```
In [22]: m = np.mean(credit_df['Age'])
print('mean:',m)
for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        titanic_df['Age'] = titanic_df['Age'].replace(i,m)
```

mean: 29.347486437613018

imputations using median

```
In [24]: m = credit_df['Age'].median()
print("median",m)
for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,m)
```

median 27.58

imputations using zero

```
In [25]: for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,0)
```

Univariate, Bivariate and Multivariate Analysis

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import math
```

```
In [3]: card_approval_df=pd.read_csv('clean_dataset.csv')
print(card_approval_df.head())
```

	Gender	Age	Debt	Married	BankCustomer	Industry	Ethnicity	\
0	1	30.83	0.000	1	1	Industrials	White	
1	0	58.67	4.460	1	1	Materials	Black	
2	0	24.50	0.500	1	1	Materials	Black	
3	1	27.83	1.540	1	1	Industrials	White	
4	1	20.17	5.625	1	1	Industrials	White	

	YearsEmployed	PriorDefault	Employed	CreditScore	DriversLicense	\
0	1.25	1	1	1	0	
1	3.04	1	1	6	0	
2	1.50	1	0	0	0	
3	3.75	1	1	5	1	
4	1.71	1	0	0	0	

	Citizen	ZipCode	Income	Approved
0	ByBirth	202	0	1
1	ByBirth	43	560	1
2	ByBirth	280	824	1
3	ByBirth	100	3	1
4	ByOtherMeans	120	0	1

```
In [4]: print(card_approval_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 690 entries, 0 to 689
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                690 non-null   int64
1   Age                   690 non-null   float64
2   Debt                  690 non-null   float64
3   Married               690 non-null   int64
4   BankCustomer          690 non-null   int64
5   Industry              690 non-null   object
6   Ethnicity             690 non-null   object
7   YearsEmployed         690 non-null   float64
8   PriorDefault          690 non-null   int64
9   Employed              690 non-null   int64
10  CreditScore           690 non-null   int64
11  DriversLicense        690 non-null   int64
12  Citizen               690 non-null   object
13  ZipCode               690 non-null   int64
14  Income                690 non-null   int64
15  Approved              690 non-null   int64
dtypes: float64(3), int64(10), object(3)
memory usage: 86.4+ KB
None
```

```
In [5]: card_approval_df.duplicated().sum()

Out[5]: 0
```


Univariate Analysis of continuous Variables

```
In [6]: card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']].describe()
```

	Age	Debt	YearsEmployed	CreditScore	Income
count	690.000000	690.000000	690.000000	690.000000	690.000000
mean	31.514116	4.758725	2.223406	2.400000	1017.385507
std	11.860245	4.978163	3.346513	4.86294	5210.102598
min	13.750000	0.000000	0.000000	0.000000	0.000000
25%	22.670000	1.000000	0.165000	0.000000	0.000000
50%	28.460000	2.750000	1.000000	0.000000	5.000000
75%	37.707500	7.207500	2.625000	3.000000	395.500000
max	80.250000	28.000000	28.500000	67.000000	100000.000000

```
In [7]: sns.histplot(card_approval_df.Age, kde=True)

Out[7]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```




Univariate Analysis of categorical Variables

```
In [8]: sns.countplot(card_approval_df.Gender)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[8]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```




```
In [9]: sns.countplot(card_approval_df.Ethnicity)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(  
<AxesSubplot:xlabel='Ethnicity', ylabel='count'>
```

Out[9]:



Bivariate analysis of continuous variable

```
In [10]: card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']].corr()
```

```
Out[10]:
```


	Age	Debt	YearsEmployed	CreditScore	Income
Age	1.000000	0.202177	0.391464	0.187327	0.018719
Debt	0.202177	1.000000	0.298902	0.271207	0.123121
YearsEmployed	0.391464	0.298902	1.000000	0.322330	0.051345
CreditScore	0.187327	0.271207	0.322330	1.000000	0.063692
Income	0.018719	0.123121	0.051345	0.063692	1.000000

```
In [12]: sns.scatterplot(card_approval_df.YearsEmployed, card_approval_df.Income)
plt.ylim(0, 20000)
```

C:\Users\gptkgf\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
(0.0, 20000.0))
```

```
Out[12]:
```



Bivariate Analysis of Categorical Variables vs Continuous Variables

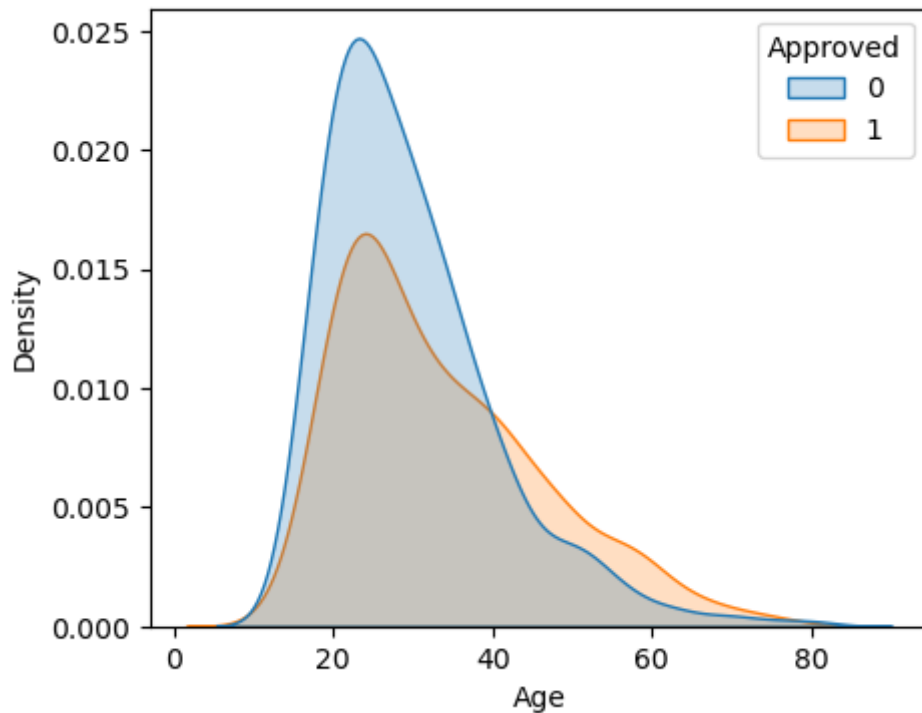
```
In [13]: card_approval_df.groupby(by='Approved').agg('mean')[['Age', 'Debt', 'YearsEmployed',
```

Out[13]:

	Age	Debt	YearsEmployed	CreditScore	Income
Approved					
0	29.773029	3.839948	1.257924	0.631854	198.605744
1	33.686221	5.904951	3.427899	4.605863	2038.859935

```
In [19]: plt.figure(figsize=(5,4))  
sns.kdeplot(data=card_approval_df,x='Age',hue='Approved',fill=True)
```


Out[19]: <AxesSubplot:xlabel='Age', ylabel='Density'>



Bivariate Analysis of Categorical Variables vs Categorical Variables

```
In [20]: sns.countplot(data=card_approval_df,x='Approved',hue='Gender')
```

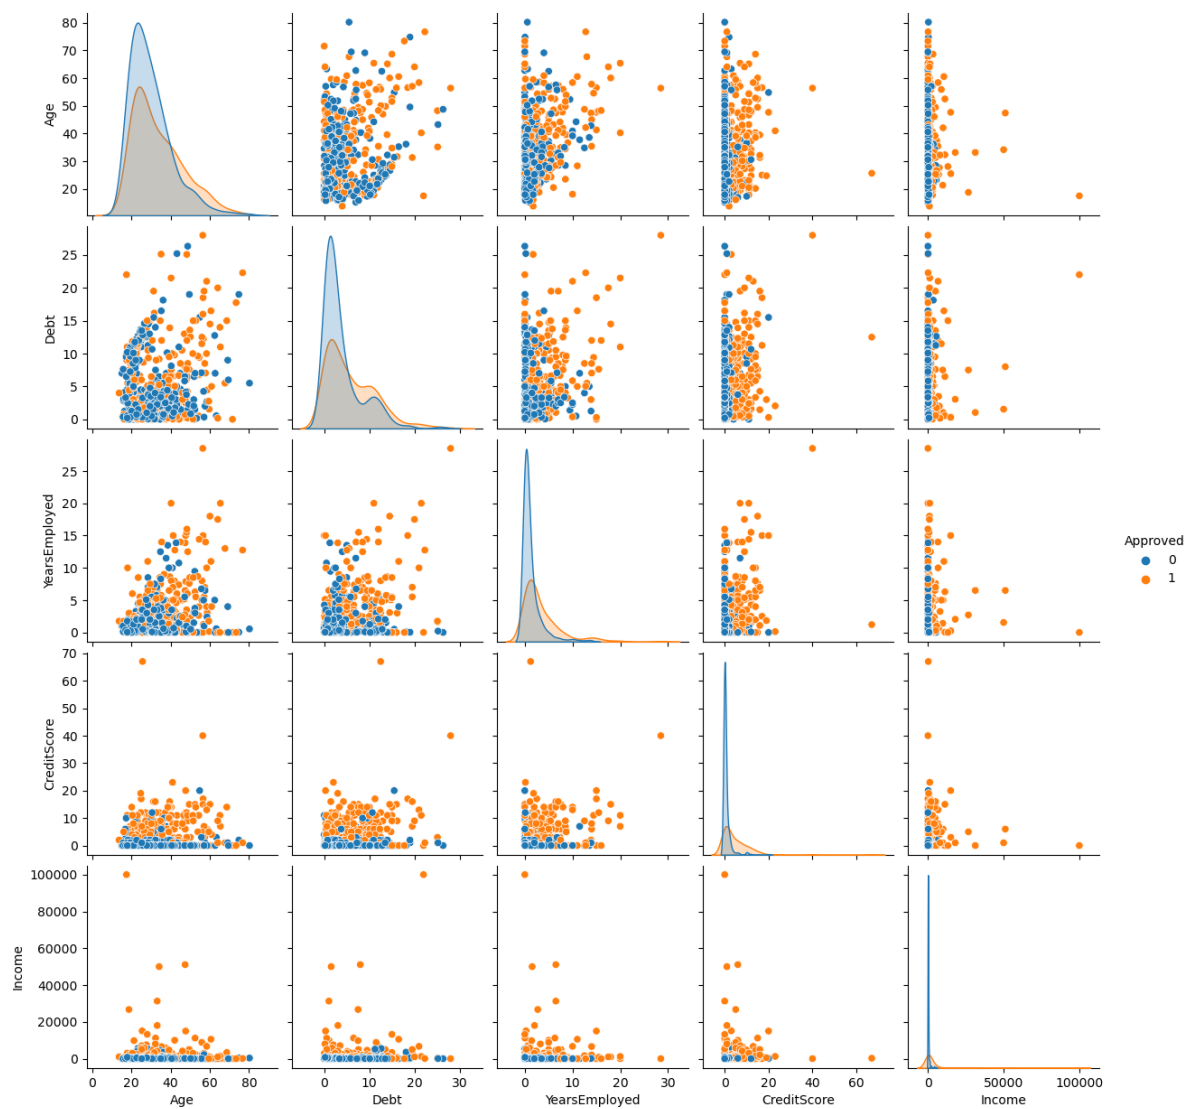
Out[20]: <AxesSubplot:xlabel='Approved', ylabel='count'>



Multivariate Analysis

```
In [21]: sns.pairplot(data=card_approval_df[['Age', 'Debt', 'YearsEmployed', 'CreditScore', 'Income']])
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x23b342f37c0>
```



In []: